

Ahmed Abdulla

[✉️ aaa384@cornell.edu](mailto:aaa384@cornell.edu) | [📞 929-312-9571](tel:929-312-9571) | [LinkedIn](#) | [GitHub](#) | [Portfolio](#)

Education

Cornell University, Bachelor of Science | **GPA: 3.7/4.0**

Expected May 2027

- **Major:** Information Science **Minor:** Data Science
- **Coursework:** Operating Systems, Systems Architecture, Databases, Data Structures & Algorithms, Machine Learning, Functional Programming, Object-Oriented Programming, Probability & Statistics, Linear Algebra, Discrete Math
- **Honors:** Dean's Honors List, 1st Place – Cornell Hackathon

Work Experience

Incoming Software Engineering Intern | Datadog | New York, NY

Summer 2026

- Distributed Systems, Metrics Team

Research Software Engineer | Cornell Lab of Ornithology | Ithaca, NY

Sep 2025 – Present

- Architected scalable data pipelines combining drone imagery and acoustic recordings using Python, FastAPI, and PostgreSQL.
- Fine-tuned deep learning models for species detection, sound separation, and density estimation.
- Built internal APIs and data services delivering near-real-time analytics to researchers through a React-based dashboard.
- Implemented monitoring and automated retraining workflows using Docker, Prefect, and Google Cloud Storage for large-scale ecological datasets.

Software Engineering Intern | Simulacrum Inc. | New York, NY

May 2025 – Aug 2025

- Launched core backend powering time-series analytics, processing 500K+ API requests/month at sub-200ms latency.
- Developed TypeScript SDK simplifying client integration and reducing setup time by 40%.
- Deployed containerized services using Docker, FastAPI, and Nginx, routing live model requests with zero downtime.
- Integrated CI/CD workflows and production monitoring using GitHub Actions, Prometheus, and Grafana.

Software Engineering Intern | Iconnectiv LLC | Bridgewater, NJ

May 2024 – Aug 2024

- Developed Python ETL pipeline to clean and load 14M+ GNS records into Oracle SQL, improving runtime efficiency by 35%.
- Built Angular/PrimeNG interface enabling three departments to query processed data in real time.
- Automated quarterly refresh and validation tasks using Bash, improving accuracy by 40%.

Leadership & Involvements

Cornell Data Strategy | Vice President, Director of Tech

Aug 2024 – Present

- Led 25+ developers across 3 pods to deliver 4+ production releases for enterprise clients at 600M+ scale.
- Standardized CI/CD and on-call rotations, reducing incident resolution time by 40%.
- Mentored junior developers and coordinated cross-pod collaboration to improve code quality and team efficiency.

Technical Projects

JobLink – AI Job Application Tracker | *scikit-learn, spaCy, RandomForest, OpenAI LLM*

Apr 2025

- Developed a full-stack platform to centralize job application emails with secure APIs and tracking dashboard.
- Built hybrid ML+rule classifier (TF-IDF, RandomForest, heuristics) to detect pipeline emails; added LLM verification for low-confidence cases.
- Engineered NLP pipeline with spaCy + LLMs achieving **97.8% precision** on entity extraction (status,dates, job IDs).

Poultry Health Monitoring Platform – 1st Place, Cornell Hackathon | *PyTorch, TensorFlow, CUDA*

Feb 2025

- Developed a multimodal AI platform combining acoustic and visual data to detect and prevent poultry diseases.
- Implemented CNN and audio-classification pipelines with GPU-optimized training, achieving 92% precision on live farm data.
- Built a real-time monitoring dashboard enabling farmers to track health alerts and environmental conditions remotely.

Publications

Goktas, D., Riaño-Briceño, G., Abdulla, A., et al. (2025). *TempusBench: An Evaluation Framework for Time-Series Forecasting*. NeurIPS 2025 Workshop on Recent Advances in Time Series Foundation Models (BERT2S). [OpenReview](#)

Technical Skills

- **Languages:** Java, Python, C/C++, JavaScript, TypeScript, SQL, OCaml, Assembly, HTML/CSS
- **Frameworks/Tools:** React, Spring Boot, Flask, FastAPI, Docker, Kubernetes, Jenkins, Jira, Redis
- **Databases/Cloud:** PostgreSQL, MySQL, DynamoDB, Firebase, AWS (S3, EC2, Lambda, CloudWatch), Google Cloud
- **Machine Learning:** PyTorch, TensorFlow, Scikit-learn, Transformers, Pandas, NumPy